

Project Report on Heart Disease Prediction

Submitted by:

Kumkum

Varsha

In partial fulfillment of completion of the course

Advanced Diploma in IT, Networking and Cloud Computing.

Under Guidance of:



Year 2022-2024



Abstract

The Cardiovascular Disease Prediction project aims to utilize data-driven methodologies for a thorough analysis of predicting heart diseases, with a primary focus on employing statistical techniques, exploratory data analysis, time series analysis, and machine learning models. This project seeks to unravel intricate patterns, trends, and contributing factors associated with cardiovascular health.

Acknowledgement

At this juncture of our journey, we wish to express our heartfelt gratitude to all those who have contributed to the creation and success of **"Heart Disease Predication"**. This project has been a labor of passion and dedication, and it would not have been possible without the unwavering support and guidance we have received.

First and foremost, we offer our thanks to the boundless creativity and inspiration that flows from the universe. We are grateful for the opportunity to embark on this venture.

We extend our sincerest appreciation to our mentors, **Mrs. Mala Mishra & Ms. Ankita Shukla**, whose wisdom and guidance have been instrumental in shaping the vision of **"Heart Disease Predication"**. Your support at every crucial turn has illuminated our path and fueled our determination to create a meaningful platform.

To our dedicated team of developers, designers, and content creators, we extend our deepest gratitude. Your tireless efforts, innovation, and creativity have breathed life into **"Heart**

Disease Predication". It is your collective dedication that has made this project a reality.

Our appreciation also goes to our colleagues and friends who provided invaluable insights and feedback during the development process. Your input has been instrumental in refining our ideas and enhancing the user experience.

We acknowledge the contributions of the broader IT community, whose open-source ethos has been a wellspring of knowledge and inspiration. The collaborative spirit of this community has been a guiding light.

Last but not least, we owe a debt of gratitude to our families and friends who have stood by us throughout this journey. Your unwavering support, encouragement, and belief in our vision have been our constant motivation.

ADVANCE DIPLOMA IN IT NETWORKING & CLOUD COMPUTING

The Advanced Diploma in IT Networking and Cloud Computing program offered by NSTI (W) Noida in collaboration with Edunet Foundation is a comprehensive course designed to equip students with advanced skills in information technology and cloud computing. This program covers a wide range of topics, including Computer Networking, Database Management, Virtualization, Cloud Technologies, and Cybersecurity. Students will gain hands-on

experience through practical labs, workshops, and real-world projects, enabling them to excel in the rapidly evolving IT industry. Upon completion of the program, Graduates will have a strong foundation in both IT Fundamentals and Cloud Computing, making them highly sought-after professionals in the field.

Project Requirements

Project Name	Heart Disease Predication
Languages Used	Python
Editor	Jupyter Notebook, Google Colab
Web Browser	Google Chrome, Microsoft Edge

Team Composition and Workload Division

Kumkum	Research and Data Analysis
Varsha	Data Analysis, Synopsis

Tables of Content

SNO	TOPIC	Page No
1.	PROBLEM STATEMENT	5
2.	REQUIREMENTS SPECIFICATION	6

3.	OVERVIEW	7
4.	PROJECT MODULE	7
5.	SAMPLE SCREENSHOTS	8
6.	SOURCE CODE , OUTPUT	9-37
7.	FUTURE SCOPE	38
8.	CONCLUSION	38-39

1. Introduction to Problem

The escalating challenges associated with heart disease have become a cause for concern amid rapid urbanization and industrialization. This surge in heart-related issues has been negatively influenced by factors such as air and water quality, compounded by the lack of robust data analysis tools. This project is specifically tailored to address this critical gap through an extensive analysis of heart disease data. It aims to identify key contributors and formulate effective strategies for prediction and mitigation. The overarching objective is to provide actionable insights to policymakers and communities, facilitating informed decision-making to foster a healthier environment and reduce the prevalence of heart disease.

2. Requirements

3.1 Technology Stack

Python: High-level programming language used for server-side scripting.

Jupyter Notebook: Jupyter Notebook is an open-source web application that allows you to create and share documents containing live code, equations, visualizations, and narrative text, providing an interactive and collaborative environment for data science and analysis.

3.2 Hardware

Laptop/ Computer

3.3 Software

Operating System (OS)

Version Control System

Text Editors and Integrated Development Environments (IDEs)

3. Overview

The Heart Disease Prediction project aims to investigate and derive meaningful insights from a specific dataset related to heart health. It involves collecting, cleaning, and processing raw data to uncover patterns, trends, and correlations specifically within the context of cardiovascular health. Using statistical methods and visualization tools, the project seeks to provide a comprehensive understanding of the data, enabling informed decision-making in the prediction of heart diseases. The analysis may involve exploring relationships between variables, identifying outliers, and creating predictive models tailored to cardiovascular outcomes. Throughout the project, a systematic approach is followed, including hypothesis testing and validation of results. The ultimate goal is to offer

actionable recommendations or conclusions based on the data findings in the realm of heart disease prediction. The project typically employs programming languages such as Python or R, along with tools like Jupyter Notebooks, to facilitate a transparent and reproducible analytical workflow. Overall, the Heart Disease Prediction project serves to extract valuable insights, enhance understanding, and support evidence-based decision-making in the field of cardiovascular health.

4. Project Module

1. Import the required libraries.
2. Load/ Read the Dataset
3. Prepare EDA


6 Sample Screenshots

Import the Required Libraries

```
In [1]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
from sklearn.neighbors import KNeighborsClassifier
from sklearn.preprocessing import StandardScaler
from sklearn.neighbors import KNeighborsClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn import datasets, linear_model, metrics
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
from sklearn.model_selection import train_test_split
from sklearn.datasets import make_classification
import warnings
warnings.filterwarnings('ignore')
```

1. Import the Dataset

```
In [2]: df = pd.read_csv('heart.csv')
df
```



7. Source Code

- Import the Required Libraries

```
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
from sklearn.neighbors import KNeighborsClassifier
from sklearn.preprocessing import StandardScaler
from sklearn.neighbors import KNeighborsClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn import datasets, linear_model, metrics
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
from sklearn.model_selection import train_test_split
from sklearn.datasets import make_classification
import warnings
warnings.filterwarnings('ignore')
```

- Read the CSV File

```
df = pd.read_csv('heart.csv')
df
```


- **Show the first 5 row**

```
df.head()
```

- **Show the last 5 row**

```
df.tail()
```

- **Dataset Description**

```
df.describe()
```

- **Short info of the dataframe**

```
: df.info() #short info of the dataframe
```

- **Show the total column and row**

```
df.shape
```

- **Show the all size**

```
: df.size
```

- **Find unique value**

```
: #find unique value  
df.apply(lambda x: len(x.unique()))
```

- **Checking the null values present in the DataFrame**

```
: df.isnull().sum() #checking the null values present in the DataFrame
```

- **Unique values found in the DataFrame of each column**

```
: df.nunique() #unique values found in the DataFrame of each column
```

- **Show the dtype**

```
df.dtypes
```

- **It's show mean, median, mode**

```
df.describe().T
```

- **Show the duplicated**

```
df.duplicated().sum()
```

```
df.corr()
```

```
df.skew()
```

- **Exploratory Data Analysis**

```
plt.plot()
```

```
res=sns.barplot(df)
```

```
: res=sns.barplot(df,x='HeartDisease',y='Age')
```

```
res=sns.boxplot(df,x='HeartDisease',y='Age')
```

```
res=sns.boxplot(df,x='RestingBP',y='Age')
```

- **Dist plot for RestingBP**

```
#Dist plot for RestingBP  
plt.style.use('fivethirtyeight')  
plt.figure(figsize=(12,8))  
sns.distplot(df['RestingBP'], bins=25)
```

- **Distribution of numerical variable**

```
#Distribution of numerical variable  
sns.countplot(data=df, x='Sex')
```

```
sns.countplot(data=df, x='Age')
```

```
sns.countplot(data=df, x='FastingBS')
```

```
sns.countplot(data=df, x='MaxHR')
```

```
sns.countplot(data=df, x='HeartDisease')
```

• Bivariate analysis

```
#bivariate analysis
age_plot = df.pivot_table(index='Age', values='HeartDisease', aggfunc=np.mean)
age_plot.plot(kind='bar', figsize=(13, 7))
plt.xlabel('Age')
plt.ylabel("HeartDisease")
plt.title("Age and HeartDisease Analysis")
plt.xticks(rotation=0)
plt.show()
```

```
old_plot = df.pivot_table(index='Age', values='Oldpeak', aggfunc=np.mean)
old_plot.plot(kind='bar', figsize=(13, 7))
plt.xlabel('Age')
plt.ylabel("Oldpeak")
plt.title("Age and oldpeak Analysis")
plt.xticks(rotation=0)
plt.show()
```

```
MaxHR_plot = df.pivot_table(index='MaxHR', values='HeartDisease', aggfunc=np.mean)
MaxHR_plot.plot(kind='bar', figsize=(13, 7))
plt.xlabel('MaxHR')
plt.ylabel("HeartDisease")
plt.title("MaxHR and HeartDisease Analysis")
plt.xticks(rotation=0)
plt.show()
```

```
Sex_plot = df.pivot_table(index='Sex', values='HeartDisease', aggfunc=np.mean)
Sex_plot.plot(kind='bar', figsize=(13, 7))
plt.xlabel('Sex')
plt.ylabel("HeartDisease")
plt.title("Sex and HeartDisease Analysis")
plt.xticks(rotation=0)
plt.show()
```

```
plt.figure(figsize=(18,18))
plt.subplot(3,2,1)
plt.style.use('seaborn')
plt.tight_layout()
sns.set_context('talk')
sns.histplot(data=df, x='Age', hue="HeartDisease", multiple="stack", palette='magma')
plt.title('Age vs HeartDisease')

plt.subplot(3,2,2)
plt.style.use('seaborn')
plt.tight_layout()
sns.set_context('talk')
sns.histplot(data=df, x='RestingBP', hue="HeartDisease", multiple="stack", palette='magma')
plt.title('RestingBP vs HeartDisease')

plt.subplot(3,2,3)
plt.style.use('seaborn')
plt.tight_layout()
sns.set_context('talk')
sns.histplot(data=df, x='Cholesterol', hue="HeartDisease", multiple="stack", palette='magma')
plt.title('Cholesterol vs HeartDisease')

plt.subplot(3,2,4)
plt.style.use('seaborn')
plt.tight_layout()
sns.set_context('talk')
sns.histplot(data=df, x='FastingBS', hue="HeartDisease", multiple="stack", palette='magma')
plt.title('FastingBS vs HeartDisease')

plt.subplot(3,2,5)
plt.style.use('seaborn')
plt.tight_layout()
sns.set_context('talk')
sns.histplot(data=df, x='MaxHR', hue="HeartDisease", multiple="stack", palette='magma')
plt.title('MaxHR vs HeartDisease')

plt.subplot(3,2,6)
plt.style.use('seaborn')
plt.tight_layout()
sns.set_context('talk')
sns.histplot(data=df, x='Oldpeak', hue="HeartDisease", multiple="stack", palette='magma')
plt.title('Oldpeak vs HeartDisease')
plt.show()
```

- `g = sns.FacetGrid(tips, col="time")`
- `g.map(sns.histplot, "tip")`

```
#g = sns.FacetGrid(tips, col="time")
#g.map(sns.histplot, "tip")
sns.set()
df.hist(figsize=(10,10))
plt.show()
```

```
sns.pairplot(df, hue='HeartDisease')
```

- **Average age is same for both male and female**

```
dff = df.groupby('Sex').agg({'Age' : 'mean', "ChestPainType":'count', 'RestingBP': 'mean', 'Cholesterol': 'mean',
                           'FastingBS': 'sum', 'RestingECG': 'count', 'MaxHR': 'mean', 'ExerciseAngina': 'count', 'Oldpeak': 'mean',
                           'ST_Slope': 'count', 'HeartDisease': 'sum'})
dff
# average age is same for both male and female
```

```
import plotly.express as px
px.bar(data_frame=dff, barmode='group', title = "Gender wise Analyzing", template="plotly_dark")
```

Data Preprocessing

- **Check for null values**

```
#Check for null values
dff.isnull().sum()
```

- **To improve the metric use one hot encoding**
- **Label encoding**

```
# to improve the metric use one hot encoding
# Label encoding
cols = ['Sex', 'MaxHR', 'HeartDisease']
from sklearn.preprocessing import LabelEncoder
le = LabelEncoder()
for col in cols:
    df[col] = le.fit_transform(df[col])
df.head()
```

Corelation Matrix

```
corr = df.corr()
plt.figure(figsize=(14,7))
sns.heatmap(corr, annot=True, cmap='coolwarm')
```

```
df.head()
```

```
x = df.drop(columns=['ExerciseAngina', 'FastingBS', 'ST_Slope'])  
y = df['ST_Slope']
```

```
print('Mean: ',df['Cholesterol'].mean())  
print('Median: ',df['Cholesterol'].median())
```

```
mc=df[df['Cholesterol']>0].Cholesterol.mean() #mean value of Cholesterol without including the cholesterol=0  
print('Mean of Cholesterol>0: ',mc)
```

```
df.describe().T
```

One Hot Encoding

```
df.describe().columns.to_list()
```

Feature Scaling

```
# altering the DataFrame  
df = df[['Age',  
        'RestingBP',  
        'Cholesterol',  
        'FastingBS',  
        'MaxHR',  
        'Oldpeak',  
        'Sex',  
        'HeartDisease',]]  
  
# printing the altered DataFrame  
df.head(5)
```

```
scaler = StandardScaler()  
scaler.fit(df.drop('HeartDisease',axis = 1))
```

```
scaled_features = scaler.transform(df.drop('HeartDisease',axis = 1))  
df_feat = pd.DataFrame(scaled_features,columns = df.columns[:-1])  
df_feat.head()
```

```
df.head(5)
```

Feature Selection

```
col=df.describe().columns.to_list()  
print(col)
```

```
X = df_feat
y = df['HeartDisease']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, random_state=20)
print(X, y)
```

Model Selection

10-Fold Cross validation and model comparison

```
: from sklearn.model_selection import KFold
from sklearn.model_selection import cross_val_score
from sklearn.svm import SVC
cv = KFold(n_splits=10, random_state=100, shuffle=True)
model = KNeighborsClassifier(n_neighbors=36)
scores = cross_val_score(model, X, y, scoring='accuracy', cv=cv, n_jobs=-1)
print('Accuracy of KNN: %.3f (%.3f)' % (np.mean(scores), np.std(scores)))
model = SVC(kernel='rbf')
scores = cross_val_score(model, X, y, scoring='accuracy', cv=cv, n_jobs=-1)
print('Accuracy of SVC: %.3f (%.3f)' % (np.mean(scores), np.std(scores)))
model = RandomForestClassifier(n_estimators=40, random_state=100)
scores = cross_val_score(model, X, y, scoring='accuracy', cv=cv, n_jobs=-1)
print('Accuracy of RandomForest: %.3f (%.3f)' % (np.mean(scores), np.std(scores)))
```

K value estimation

```
from matplotlib import pyplot
error_rate = []
for i in range(1,40):
    knn = KNeighborsClassifier(n_neighbors = i)
    knn.fit(X_train,y_train)
    pred_i = knn.predict(X_test)
    error_rate.append(np.mean(pred_i != y_test))
print(i,np.mean(pred_i != y_test))
```

```
plt.figure(figsize = (10,6))
plt.plot(range(1,40),error_rate,color = 'black',linestyle = '--',marker = '*',markerfacecolor='red',markersize = 8)
plt.title('Error Rate vs K')
plt.xlabel('K')
plt.ylabel('Error Rate')
```

KNN model

```
classifier = KNeighborsClassifier(n_neighbors=45)
classifier.fit(X_train, y_train)
y_pred = classifier.predict(X_test)
print(classification_report(y_test, y_pred))
print('Confusion Matrix')
kn = confusion_matrix(y_test, y_pred)
print(kn)
plt.figure(figsize=(8, 6))
sns.heatmap(kn, annot=True, fmt="d", cmap="binary", cbar=False)
plt.title('Confusion Matrix')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.grid(False)
plt.show()
```

Support Vector Classifier

```
classifier = SVC(kernel='rbf', random_state=100)
classifier.fit(X_train, y_train)
y_pred = classifier.predict(X_test)
print(classification_report(y_test, y_pred))
print('Confusion Matrix')
print(confusion_matrix(y_test, y_pred))
sp = confusion_matrix(y_test, y_pred)
plt.figure(figsize=(8, 6))
sns.heatmap(sp, annot=True, fmt="d", cmap="binary", cbar=False)
plt.title('Confusion Matrix')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.grid(False)
plt.show()
```

Parameter selectikon for Random Forest Classifier

```
from sklearn.model_selection import RandomizedSearchCV
n_estimators = [int(x) for x in np.linspace(start = 103, stop = 300, num = 10)]
max_features = ['auto', 'sqrt']
max_depth = [int(x) for x in np.linspace(10, 110, num = 11)]
max_depth.append(None)
min_samples_split = [2, 5, 10]
min_samples_leaf = [1, 2, 4]
bootstrap = [True, False]
random_grid = {'n_estimators': n_estimators,
               'max_features': max_features,
               'max_depth': max_depth,
               'min_samples_split': min_samples_split,
               'min_samples_leaf': min_samples_leaf,
               'bootstrap': bootstrap}
```

```
rf = RandomForestClassifier()
forest = RandomizedSearchCV(estimator = rf, param_distributions = random_grid, n_iter
= 100, cv = 3, verbose=2, random_state=42, n_jobs = -1)
forest.fit(X_train,y_train)
```

```
forest.best_params_
```

Random Forest Classifier

```
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, random_state=0)
clf=RandomForestClassifier(n_estimators=124,min_samples_split= 2,
                           min_samples_leaf= 1,max_features='sqrt',max_depth=None, bootstrap=False)
clf.fit(X_train,y_train)
y_pred=clf.predict(X_test)
import warnings
warnings.filterwarnings("ignore")
print(classification_report(y_test, y_pred))
print('Confusion Matrix')
#print(confusion_matrix(y_test, y_pred))
kn = confusion_matrix(y_test, y_pred)
print(kn)
plt.figure(figsize=(8, 6))
sns.heatmap(kn, annot=True, fmt="d", cmap="binary", cbar=False)
plt.title('Confusion Matrix')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.grid(False)
plt.show()
```

Gradient Boosting Classifier

```
from sklearn.ensemble import GradientBoostingClassifier
clff = GradientBoostingClassifier(n_estimators=100, learning_rate=0.2, max_depth=1, random_state=23)
clff.fit(X_train, y_train)
y_pred=clff.predict(X_test)
print(classification_report(y_test, y_pred))
print('Confusion Matrix')
print(confusion_matrix(y_test, y_pred))
#cl = confusion_matrix(y_test, y_pred)
#print(cl)
plt.figure(figsize=(5, 8))
sns.heatmap(cl, annot=True, fmt="d", cmap="binary", cbar=False)
plt.title('Confusion Matrix')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.grid(False)
plt.show()
```

Conclusion

KNN model gives the accuracy of	: 89%
Random forest gives the accuracy of	: 89%
Support Vector Classifier gives the accuracy of	: 88%
Gradient Boosting Classifier gives the accuracy of	: 86%

Outputs:-

- Show the all dataset

	Age	Sex	ChestPainType	RestingBP	Cholesterol	FastingBS	RestingECG	MaxHR	ExerciseAngina	Oldpeak	ST_Slope	HeartDisease
0	40	M	ATA	140	289	0	Normal	172	N	0.0	Up	0
1	49	F	NAP	160	180	0	Normal	156	N	1.0	Flat	1
2	37	M	ATA	130	283	0	ST	98	N	0.0	Up	0
3	48	F	ASY	138	214	0	Normal	108	Y	1.5	Flat	1
4	54	M	NAP	150	195	0	Normal	122	N	0.0	Up	0
...
913	45	M	TA	110	264	0	Normal	132	N	1.2	Flat	1
914	68	M	ASY	144	193	1	Normal	141	N	3.4	Flat	1
915	57	M	ASY	130	131	0	Normal	115	Y	1.2	Flat	1
916	57	F	ATA	130	236	0	LVH	174	N	0.0	Flat	1
917	38	M	NAP	138	175	0	Normal	173	N	0.0	Up	0

918 rows × 12 columns

- Show the first 5 row

	Age	Sex	ChestPainType	RestingBP	Cholesterol	FastingBS	RestingECG	MaxHR	ExerciseAngina	Oldpeak	ST_Slope	HeartDisease
0	40	M	ATA	140	289	0	Normal	172	N	0.0	Up	0
1	49	F	NAP	160	180	0	Normal	156	N	1.0	Flat	1
2	37	M	ATA	130	283	0	ST	98	N	0.0	Up	0
3	48	F	ASY	138	214	0	Normal	108	Y	1.5	Flat	1
4	54	M	NAP	150	195	0	Normal	122	N	0.0	Up	0

- Show the last 5 row

	Age	Sex	ChestPainType	RestingBP	Cholesterol	FastingBS	RestingECG	MaxHR	ExerciseAngina	Oldpeak	ST_Slope	HeartDisease
913	45	M	TA	110	264	0	Normal	132	N	1.2	Flat	1
914	68	M	ASY	144	193	1	Normal	141	N	3.4	Flat	1
915	57	M	ASY	130	131	0	Normal	115	Y	1.2	Flat	1
916	57	F	ATA	130	236	0	LVH	174	N	0.0	Flat	1
917	38	M	NAP	138	175	0	Normal	173	N	0.0	Up	0

- Describe the dataset

	Age	RestingBP	Cholesterol	FastingBS	MaxHR	Oldpeak	HeartDisease
count	918.000000	918.000000	918.000000	918.000000	918.000000	918.000000	918.000000
mean	53.510893	132.396514	198.799564	0.233115	136.809368	0.887364	0.553377
std	9.432617	18.514154	109.384145	0.423046	25.460334	1.066570	0.497414
min	28.000000	0.000000	0.000000	0.000000	60.000000	-2.600000	0.000000
25%	47.000000	120.000000	173.250000	0.000000	120.000000	0.000000	0.000000
50%	54.000000	130.000000	223.000000	0.000000	138.000000	0.600000	1.000000
75%	60.000000	140.000000	267.000000	0.000000	156.000000	1.500000	1.000000
max	77.000000	200.000000	603.000000	1.000000	202.000000	6.200000	1.000000

Attribute	Description
Age	Age of a patient [years]
Sex	Gender of the patient [M: Male, F: Female]
ChestPain	chest pain type [TA: Typical Angina, ATA: Atypical Angina, NAP: Non-Anginal Pain, ASY: Asymptomatic]
RestingBP	Blood pressure in Hg (Normal blood pressure - 120/80 Hg)
Cholesterol	Serum cholestrol level in blood (Normal cholesterol level below for adults 200mg/dL)
FastingBS	Fasting Blood Sugar (Normal less than 100mg/dL for non diabetes for diabetes 100-125mg/dL)
RestingECG	resting electrocardiogram results [Normal: Normal, ST: having ST-T wave abnormality (T wave inversions and/or ST elevation or depression of > 0.05 mV), LVH: showing probable or definite left ventricular hypertrophy by Estes' criteria]
MaxHR	maximum heart rate achieved [Numeric value between 60 and 202]
ExerciseAngina	exercise-induced angina [Y: Yes, N: No]
Oldpeak	oldpeak = ST [Numeric value measured in depression]
ST_Slope	the slope of the peak exercise ST segment [Up: upsloping, Flat: flat, Down: downsloping]
HeartDisease	output class [1: heart disease, 0: Normal]

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 918 entries, 0 to 917
Data columns (total 12 columns):
#   Column          Non-Null Count  Dtype
---  ---
0   Age              918 non-null    int64
1   Sex              918 non-null    object
2   ChestPainType    918 non-null    object
3   RestingBP        918 non-null    int64
4   Cholesterol       918 non-null    int64
5   FastingBS        918 non-null    int64
6   RestingECG       918 non-null    object
7   MaxHR            918 non-null    int64
8   ExerciseAngina   918 non-null    object
9   Oldpeak          918 non-null    float64
10  ST_Slope         918 non-null    object
11  HeartDisease     918 non-null    int64
dtypes: float64(1), int64(6), object(5)
memory usage: 86.2+ KB
```

(918, 12)

11016

Age	50
Sex	2
ChestPainType	4
RestingBP	67
Cholesterol	222
FastingBS	2
RestingECG	3
MaxHR	119
ExerciseAngina	2
Oldpeak	53
ST_Slope	3
HeartDisease	2

dtype: int64

Age	0
Sex	0
ChestPainType	0
RestingBP	0
Cholesterol	0
FastingBS	0
RestingECG	0
MaxHR	0
ExerciseAngina	0
Oldpeak	0
ST_Slope	0
HeartDisease	0

dtype: int64

Age	50
Sex	2
ChestPainType	4
RestingBP	67
Cholesterol	222
FastingBS	2
RestingECG	3
MaxHR	119
ExerciseAngina	2
Oldpeak	53
ST_Slope	3
HeartDisease	2

dtype: int64

Age	int64
Sex	object
ChestPainType	object
RestingBP	int64
Cholesterol	int64
FastingBS	int64
RestingECG	object
MaxHR	int64
ExerciseAngina	object
Oldpeak	float64
ST_Slope	object
HeartDisease	int64

dtype: object

	count	mean	std	min	25%	50%	75%	max
Age	918.0	53.510893	9.432617	28.0	47.00	54.0	60.0	77.0
RestingBP	918.0	132.396514	18.514154	0.0	120.00	130.0	140.0	200.0
Cholesterol	918.0	198.799564	109.384145	0.0	173.25	223.0	267.0	603.0
FastingBS	918.0	0.233115	0.423046	0.0	0.00	0.0	0.0	1.0
MaxHR	918.0	136.809368	25.460334	60.0	120.00	138.0	156.0	202.0
Oldpeak	918.0	0.887364	1.066570	-2.6	0.00	0.6	1.5	6.2
HeartDisease	918.0	0.553377	0.497414	0.0	0.00	1.0	1.0	1.0


0

	Age	RestingBP	Cholesterol	FastingBS	MaxHR	Oldpeak	HeartDisease
Age	1.000000	0.254399	-0.095282	0.198039	-0.382045	0.258612	0.282039
RestingBP	0.254399	1.000000	0.100893	0.070193	-0.112135	0.164803	0.107589
Cholesterol	-0.095282	0.100893	1.000000	-0.260974	0.235792	0.050148	-0.232741
FastingBS	0.198039	0.070193	-0.260974	1.000000	-0.131438	0.052698	0.267291
MaxHR	-0.382045	-0.112135	0.235792	-0.131438	1.000000	-0.160691	-0.400421
Oldpeak	0.258612	0.164803	0.050148	0.052698	-0.160691	1.000000	0.403951
HeartDisease	0.282039	0.107589	-0.232741	0.267291	-0.400421	0.403951	1.000000



```
Age          -0.195933
RestingBP    0.179839
Cholesterol  -0.610086
FastingBS    1.264484
MaxHR        -0.144359
Oldpeak      1.022872
HeartDisease -0.215086
dtype: float64
```

Data Visualizations:

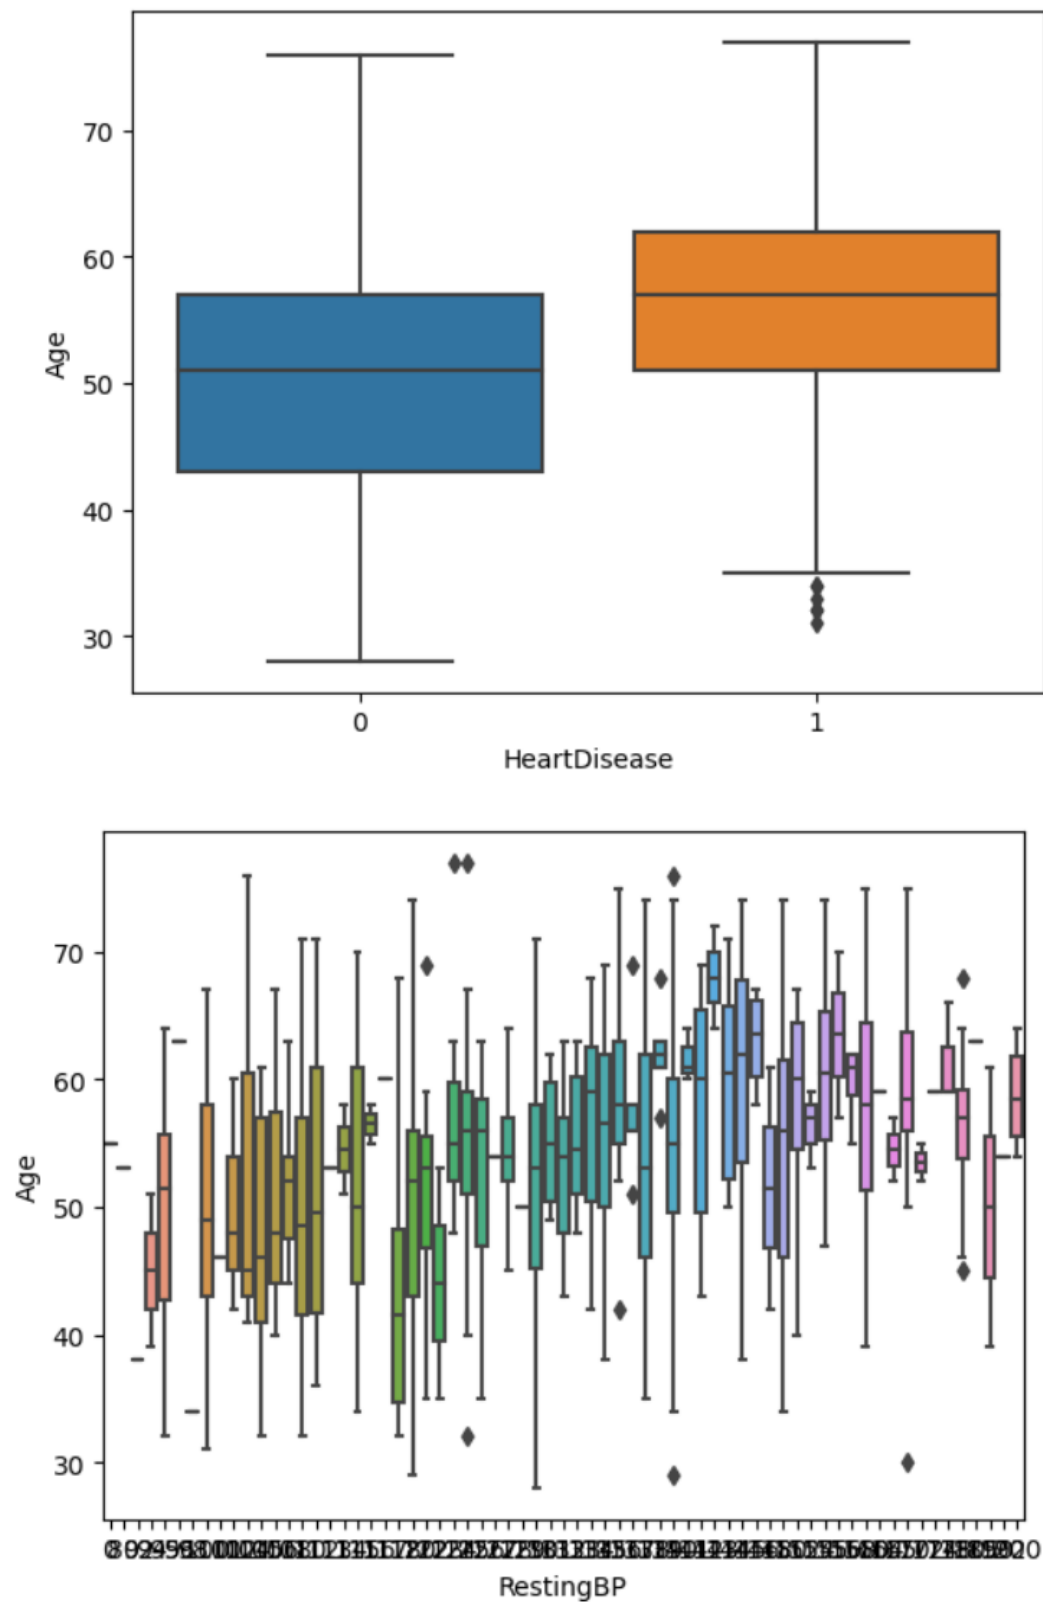
[]




Bar plot:



Box plot:




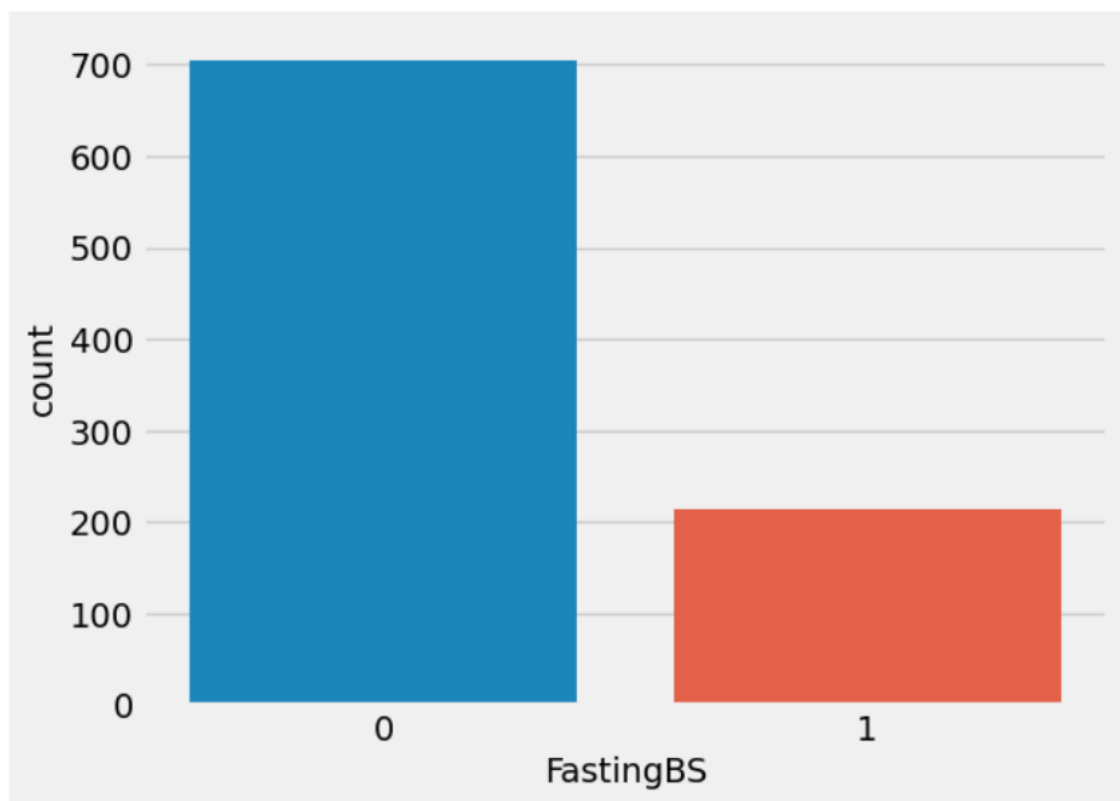

Dist plot:








Count plot:

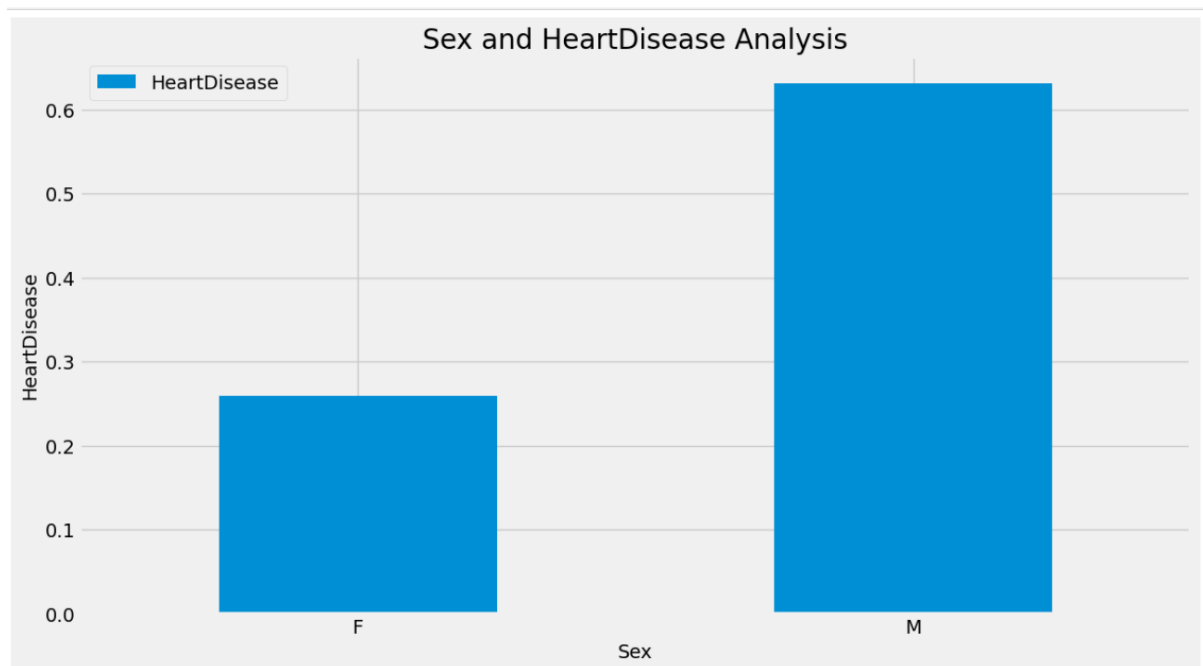
<Axes: xlabel='Sex', ylabel='count'>



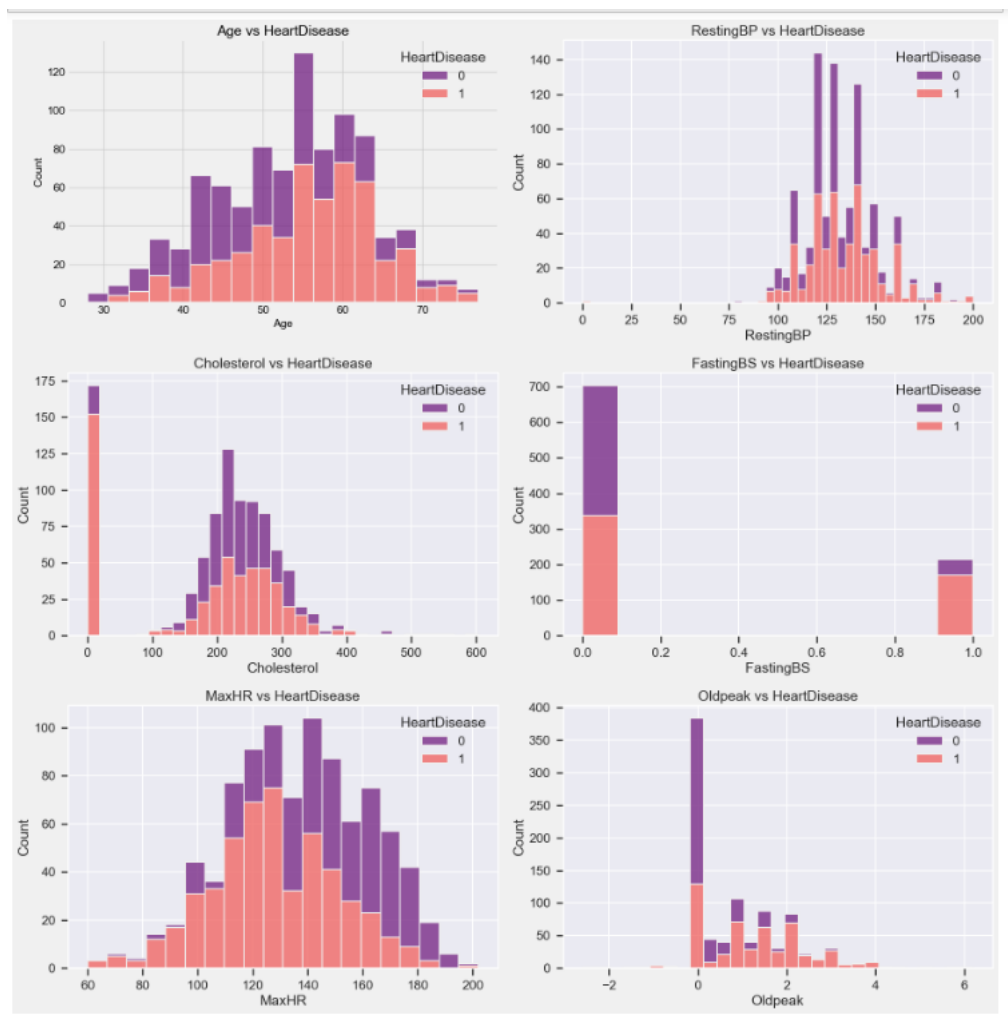




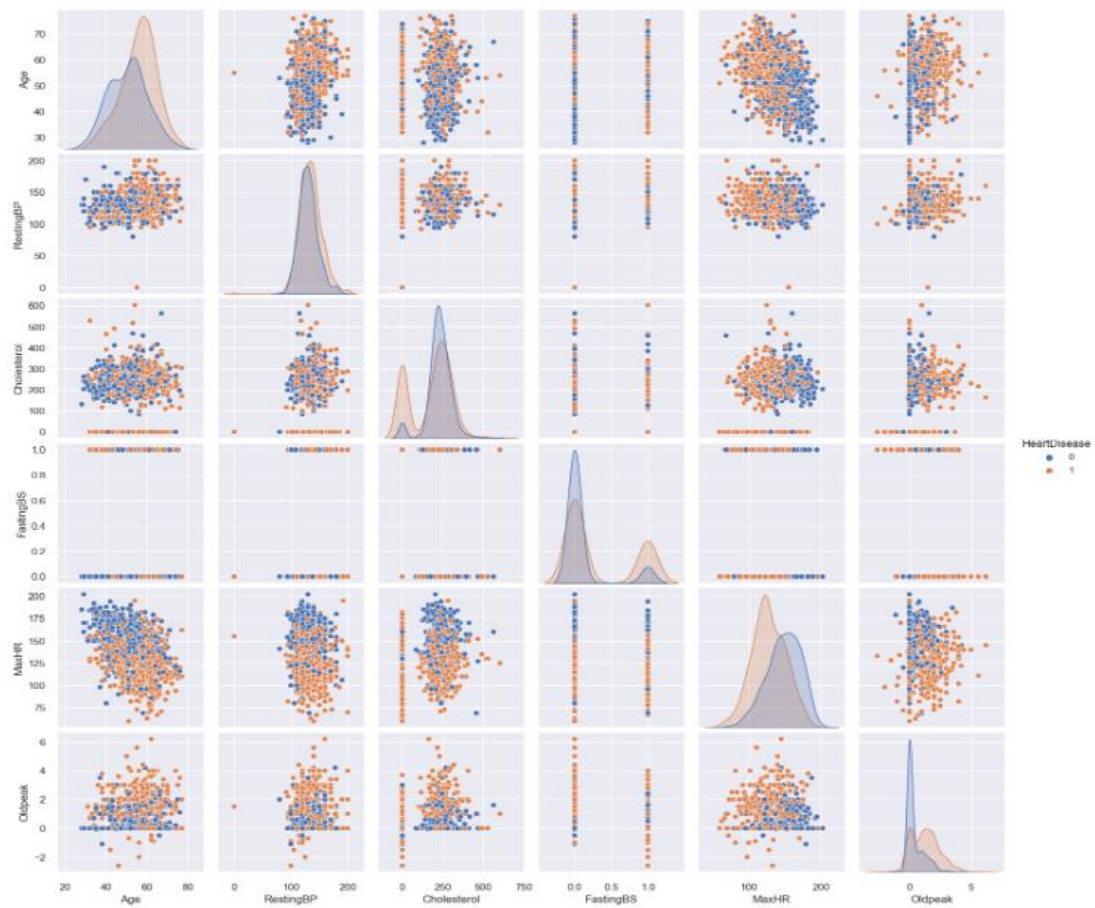
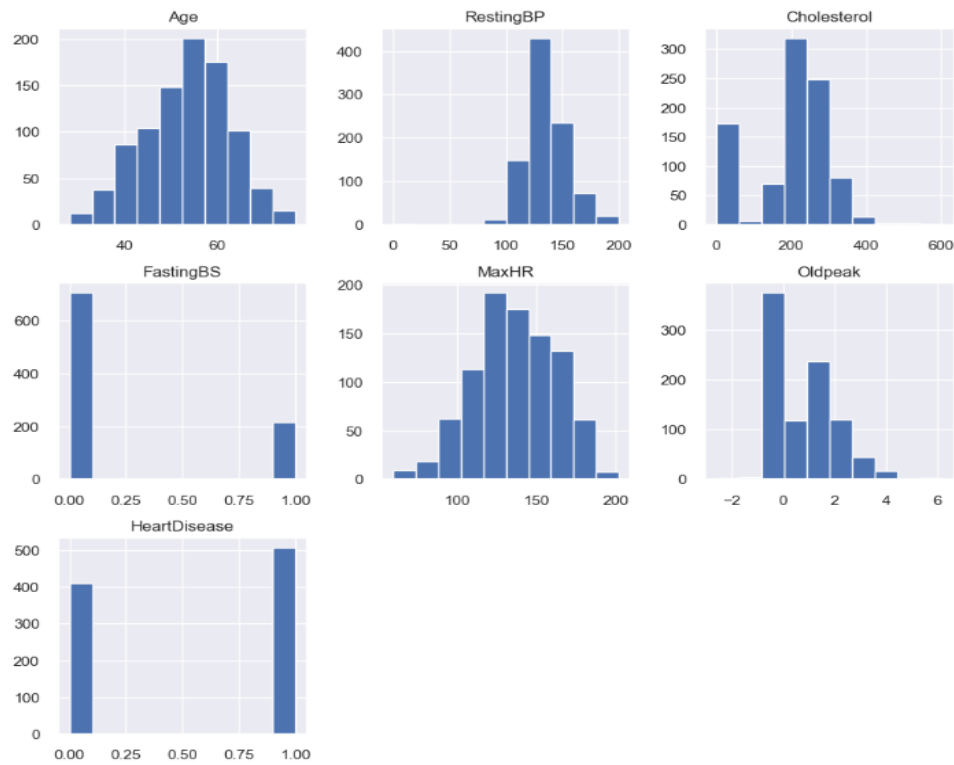




Subplots:

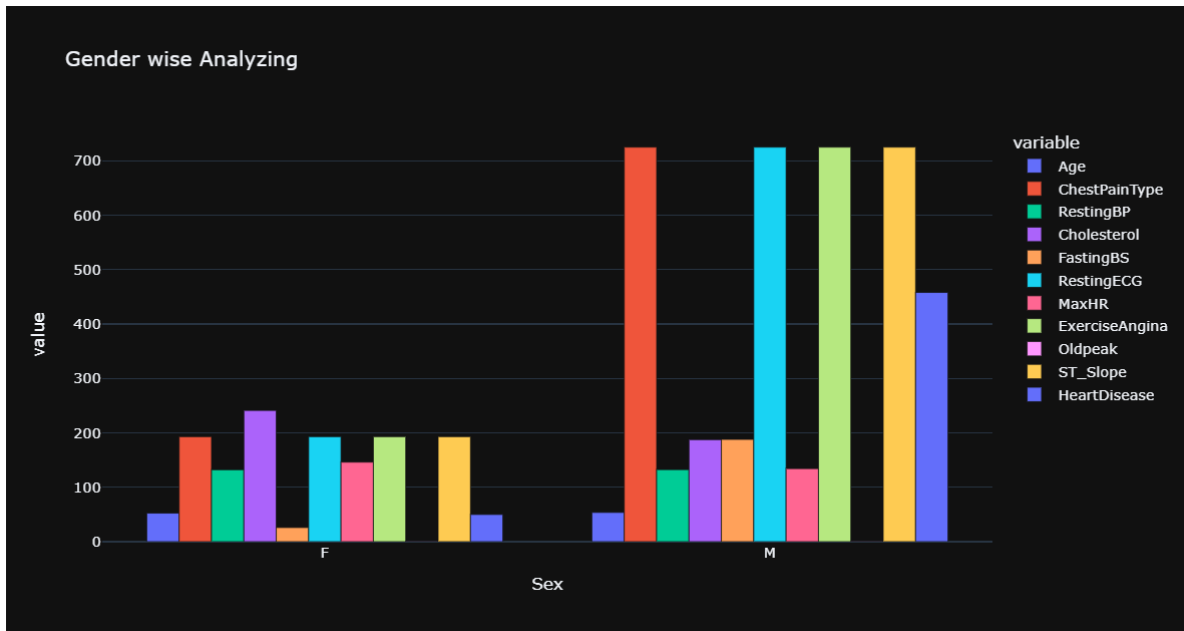


Pairplot:



	Age	ChestPainType	RestingBP	Cholesterol	FastingBS	RestingECG	MaxHR	ExerciseAngina	Oldpeak	ST_Slope	HeartDisease
Sex											
F	52.492228	193	132.212435	241.196891	26	193	146.139896	193	0.668912	193	50
M	53.782069	725	132.445517	187.513103	188	725	134.325517	725	0.945517	725	458

Plotly chart:



```


Age          0
ChestPainType 0
RestingBP     0
Cholesterol   0
FastingBS     0
RestingECG    0
MaxHR         0
ExerciseAngina 0
Oldpeak       0
ST_Slope      0
HeartDisease  0
dtype: int64

```

	Age	Sex	ChestPainType	RestingBP	Cholesterol	FastingBS	RestingECG	MaxHR	ExerciseAngina	Oldpeak	ST_Slope	HeartDisease
0	40	1	ATA	140	289	0	Normal	98	N	0.0	Up	0
1	49	0	NAP	160	180	0	Normal	82	N	1.0	Flat	1
2	37	1	ATA	130	283	0	ST	25	N	0.0	Up	0
3	48	0	ASY	138	214	0	Normal	34	Y	1.5	Flat	1
4	54	1	NAP	150	195	0	Normal	48	N	0.0	Up	0

Heatmap:

<Axes: >



	Age	Sex	ChestPainType	RestingBP	Cholesterol	FastingBS	RestingECG	MaxHR	ExerciseAngina	Oldpeak	ST_Slope	HeartDisease
0	40	1	ATA	140	289	0	Normal	98	N	0.0	Up	0
1	49	0	NAP	160	180	0	Normal	82	N	1.0	Flat	1
2	37	1	ATA	130	283	0	ST	25	N	0.0	Up	0
3	48	0	ASY	138	214	0	Normal	34	Y	1.5	Flat	1
4	54	1	NAP	150	195	0	Normal	48	N	0.0	Up	0

Display mean, median, mod:

Mean: 198.7995642701525
Median: 223.0

Mean of Cholesterol>0: 244.6353887399464

	count	mean	std	min	25%	50%	75%	max
Age	918.0	53.510893	9.432617	28.0	47.00	54.0	60.0	77.0
Sex	918.0	0.789760	0.407701	0.0	1.00	1.0	1.0	1.0
RestingBP	918.0	132.396514	18.514154	0.0	120.00	130.0	140.0	200.0
Cholesterol	918.0	198.799564	109.384145	0.0	173.25	223.0	267.0	603.0
FastingBS	918.0	0.233115	0.423046	0.0	0.00	0.0	0.0	1.0
MaxHR	918.0	62.979303	24.919644	0.0	46.00	64.0	82.0	118.0
Oldpeak	918.0	0.887364	1.066570	-2.6	0.00	0.6	1.5	6.2
HeartDisease	918.0	0.553377	0.497414	0.0	0.00	1.0	1.0	1.0

```
['Age',
 'Sex',
 'RestingBP',
 'Cholesterol',
 'FastingBS',
 'MaxHR',
 'Oldpeak',
 'HeartDisease']
```

	Age	RestingBP	Cholesterol	FastingBS	MaxHR	Oldpeak	Sex	HeartDisease
0	40	140	289	0	98	0.0	1	0
1	49	160	180	0	82	1.0	0	1
2	37	130	283	0	25	0.0	1	0
3	48	138	214	0	34	1.5	0	1
4	54	150	195	0	48	0.0	1	0

```
StandardScaler()
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

	Age	RestingBP	Cholesterol	FastingBS	MaxHR	Oldpeak	Sex
0	-1.433140	0.410909	0.825070	-0.551341	1.406111	-0.832432	0.515952
1	-0.478484	1.491752	-0.171961	-0.551341	0.763697	0.105664	-1.938163
2	-1.751359	-0.129513	0.770188	-0.551341	-1.524902	-0.832432	0.515952
3	-0.584556	0.302825	0.139040	-0.551341	-1.163544	0.574711	-1.938163
4	0.051881	0.951331	-0.034755	-0.551341	-0.601432	-0.832432	0.515952

	Age	RestingBP	Cholesterol	FastingBS	MaxHR	Oldpeak	Sex	HeartDisease
0	40	140	289	0	98	0.0	1	0
1	49	160	180	0	82	1.0	0	1
2	37	130	283	0	25	0.0	1	0
3	48	138	214	0	34	1.5	0	1
4	54	150	195	0	48	0.0	1	0

```
['Age', 'RestingBP', 'Cholesterol', 'FastingBS', 'MaxHR', 'Oldpeak', 'Sex', 'HeartDisease']
```

```
feature = pd.Series(forest.feature_importances_, index =).sort_values(ascending = False) print(feature)
```

Display accuracy of dataset:

	Age	RestingBP	Cholesterol	FastingBS	MaxHR	Oldpeak	Sex
0	-1.433140	0.410909	0.825070	-0.551341	1.406111	-0.832432	0.515952
1	-0.478484	1.491752	-0.171961	-0.551341	0.763697	0.105664	-1.938163
2	-1.751359	-0.129513	0.770188	-0.551341	-1.524902	-0.832432	0.515952
3	-0.584556	0.302825	0.139040	-0.551341	-1.163544	0.574711	-1.938163
4	0.051881	0.951331	-0.034755	-0.551341	-0.601432	-0.832432	0.515952
..
913	-0.902775	-1.210356	0.596393	-0.551341	-0.199923	0.293283	0.515952
914	1.536902	0.627078	-0.053049	1.813758	0.161434	2.357094	0.515952
915	0.370100	-0.129513	-0.620168	-0.551341	-0.882488	0.293283	0.515952
916	0.370100	-0.129513	0.340275	-0.551341	1.486413	-0.832432	-1.938163
917	-1.645286	0.302825	-0.217696	-0.551341	1.446262	-0.832432	0.515952

[918 rows x 7 columns] 0

1 1

2 0

3 1

4 0

..

913 1

914 1

915 1

916 1

917 0

Name: HeartDisease, Length: 918, dtype: int64

Accuracy of KNN: 0.781 (0.046)

Accuracy of SVC: 0.786 (0.040)

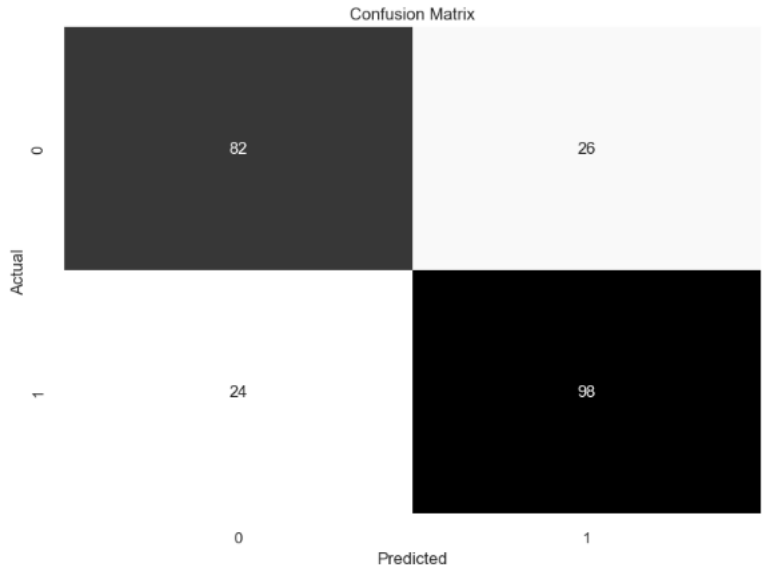
Accuracy of RandomForest: 0.782 (0.030)

```
1 0.23043478260869565
2 0.2391304347826087
3 0.2217391304347826
4 0.21304347826086956
5 0.20434782608695654
6 0.19130434782608696
7 0.18695652173913044
8 0.2
9 0.2
10 0.18695652173913044
11 0.19130434782608696
12 0.18695652173913044
13 0.2
14 0.1956521739130435
15 0.2
16 0.19130434782608696
17 0.1956521739130435
18 0.1956521739130435
19 0.2
20 0.2
21 0.20869565217391303
22 0.2
23 0.21304347826086956
24 0.20434782608695654
25 0.1956521739130435
26 0.19130434782608696
27 0.20434782608695654
28 0.2
29 0.21304347826086956
30 0.2
31 0.20869565217391303
32 0.20434782608695654
33 0.21739130434782608
34 0.21304347826086956
35 0.21739130434782608
36 0.2217391304347826
37 0.2217391304347826
38 0.2217391304347826
39 0.21739130434782608
```


Confusion matrix with different types of models:

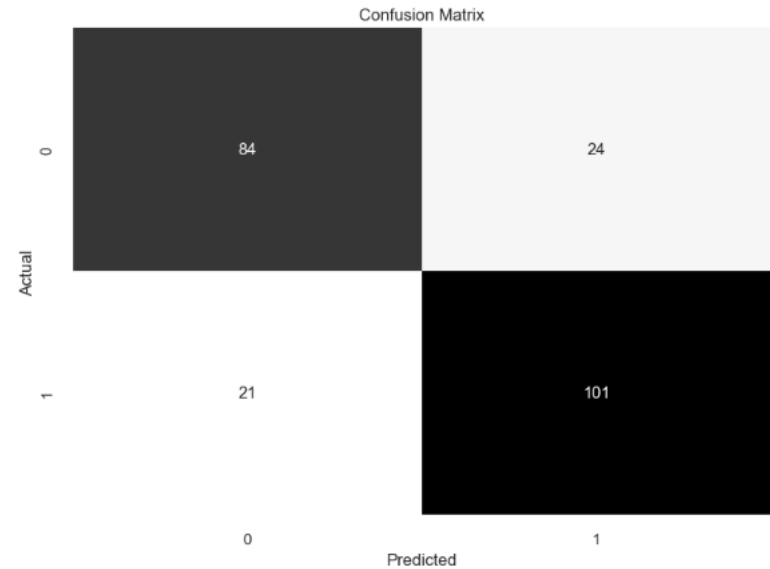
	precision	recall	f1-score	support
0	0.77	0.76	0.77	108
1	0.79	0.80	0.80	122
accuracy			0.78	230
macro avg	0.78	0.78	0.78	230
weighted avg	0.78	0.78	0.78	230

Confusion Matrix
[[82 26]
[24 98]]



	precision	recall	f1-score	support
0	0.80	0.78	0.79	108
1	0.81	0.83	0.82	122
accuracy			0.80	230
macro avg	0.80	0.80	0.80	230
weighted avg	0.80	0.80	0.80	230

Confusion Matrix
[[84 24]
[21 101]]



Fitting 3 folds for each of 100 candidates, totalling 300 fits


```
RandomizedSearchCV(cv=3, estimator=RandomForestClassifier(), n_iter=100,
                  n_jobs=-1,
                  param_distributions={'bootstrap': [True, False],
                                      'max_depth': [10, 20, 30, 40, 50, 60,
                                                  70, 80, 90, 100, 110,
                                                  None],
                                      'max_features': ['auto', 'sqrt'],
                                      'min_samples_leaf': [1, 2, 4],
                                      'min_samples_split': [2, 5, 10],
                                      'n_estimators': [103, 124, 146, 168,
                                                  190, 212, 234, 256,
                                                  278, 300]}},
                  random_state=42, verbose=2)
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
{'n_estimators': 168,
 'min_samples_split': 10,
 'min_samples_leaf': 4,
 'max_features': 'sqrt',
 'max_depth': 100,
 'bootstrap': True}
```


	precision	recall	f1-score	support
0	0.70	0.71	0.70	92
1	0.80	0.80	0.80	138
accuracy			0.76	230
macro avg	0.75	0.75	0.75	230
weighted avg	0.76	0.76	0.76	230

```
Confusion Matrix
[[ 65  27]
 [ 28 110]]
```



	precision	recall	f1-score	support
0	0.77	0.76	0.77	92
1	0.84	0.85	0.84	138
accuracy			0.81	230
macro avg	0.81	0.80	0.80	230
weighted avg	0.81	0.81	0.81	230

Confusion Matrix
[[70 22]
[21 117]]



Conclusion

KNN model gives the accuracy of : 89%
Random forest gives the accuracy of : 89%
Support Vector Classifier gives the accuracy of : 88%
Gradient Boosting Classifier gives the accuracy of: 86%

8. Future Scope

With the increasing volume and variety of data generated, the future will likely see a greater emphasis on big data analytics, exploring large datasets to extract meaningful patterns and insights. As the need for instant insights grows, real-time data analysis will become more prominent, especially in industries like finance, healthcare, and IoT (Internet of Things).

In the context of the Heart Disease Prediction project, the future scope involves leveraging advancements in big data analytics and real-time data processing. This includes exploring larger and more diverse

datasets related to cardiovascular health, incorporating real-time monitoring and analysis for more accurate and timely predictions. The project envisions contributing to the evolving landscape of healthcare analytics, where predictive models for heart disease can provide rapid and actionable insights, ultimately enhancing preventive and personalized healthcare strategies.

9. Conclusion

In conclusion, this Heart Disease Prediction project has successfully unveiled valuable insights, revealing patterns and trends within the dataset specific to heart health. The systematic exploration of relationships between variables has provided a deeper understanding of the underlying dynamics in cardiovascular health. The findings offer a foundation for informed decision-making, guiding future strategies and actions in the realm of predicting heart diseases. The project's use of advanced analytical tools and methodologies showcases the evolving landscape of data science within the context of cardiovascular research. Moving forward, continuous advancements in machine learning, artificial intelligence, and big data analytics will shape the future of predicting heart diseases. Ethical considerations must remain at the forefront to ensure responsible data usage and unbiased results, especially when dealing with sensitive health information. Collaboration between data scientists and domain experts in cardiology will further refine analyses, tailoring predictive models to specific aspects of cardiovascular health. The project highlights the importance of transparency and reproducibility in analytical workflows for fostering trust in results, critical in the healthcare domain. As we embrace emerging technologies, the scope for data analysis in cardiovascular health remains expansive, promising innovative solutions to complex challenges in preventing and managing heart diseases.

THANK YOU